

# CS 351


## Design of Large Programs

### Concurrency

October 18, 2021

# Sequential Process Characterization

- Program code (fixed)
- Control state (program counter)
- Memory state
  - stack
  - heap
- Formal properties
  - safety (does nothing wrong)
  - liveness (makes progress)




executing process

# Physical Parallelism

Parallel execution of multiple independent processes takes place on separate physical hardware resources

- multiple cores
- specialized hardware interfaces
- parallel computers
- etc.



# Logical Parallelism

- Interleaved execution of multiple independent processes takes place on a shared physical hardware resource (single CPU)
- Logical and physical parallelism coexist on modern computers
- Same two programs
  - may share a core at some point (interleaved execution)
  - may execute on separate cores at other times (parallel execution)

# Process Scheduling

- It is the responsibility of the operating system to schedule the execution of the processes sharing one computing platform
- The scheduling policy significantly impacts the execution times of the individual processes
- Any attempt to perform a performance analysis needs to take the scheduling policy into account

# Sample Scheduling Policies

- Fixed window
  - within a fixed-size window, each process has an assigned execution slot
- Round robin
  - each process gets a turn with no process being allowed to run forever
- Priority based
  - the process with the highest priority is scheduled first and runs to completion
  - the schedule may be preemptive or not

# Preemptive vs Non-Preemptive

Scenario: You are playing Candy Crush on your phone when your phone receives a phone call!

- How would you want your phone to behave?
- Would you want it to wait for you to finish playing your game?
- Or would you like it to interrupt your game and give CPU time to handling the phone call?

Probably the latter! This is the difference between preemptive and non-preemptive!

## Preemptive vs Non-Preemptive Cont.

- Preemptive: When a higher priority process comes along it is immediately given priority
- Non-Preemptive: Currently running process is allowed to finish before higher priority process is given time



# Concurrency

- *Concurrency* is an abstract unifying framework that enables one to reason about logical and physical parallelism
- It abstracts out
  - physical resources
  - timing considerations
- It achieves this by reducing all forms of parallelism to *nondeterministic execution of concurrent processes*
- It allows one to reason about the execution of concurrent processes while ignoring many of the complexities of the execution environment

# Concurrency vs Parallelism

- Concurrency
  - When two or more tasks can start, run and complete in overlapping time periods
  - Does not necessarily mean they will ever be running at the exact same time
- Parallelism
  - When tasks run at the exact same time

# Why Abstraction is Important

- Concurrent execution of multiple processes is an essential feature of modern computing
- Programming language development did not pay sufficient attention to concurrency, making programming more complex than it ought to be
- Some languages (including Java) include explicit constructs that address concurrent programming

# Why Abstraction is Important

Concurrency introduces significant levels of complexity

- programs are rarely independent of each other
- programs need to coordinate with each other and compete for resources
- programs may need to coordinate even when
  - developed independently
  - residing on processors across the world

# Fundamental Concepts

## Atomicity

- An operation is atomic if it appears to be instantaneous and uninterruptable
- Programming languages provide only minimal atomicity guarantees
  - read a primitive variable
  - write a primitive variable
  - long and double read/writes are not guaranteed to be atomic
- This greatly complicates the programming task

# Fundamental Concepts

## Fairness

- Nondeterminism abstracts out the details of the scheduling policy
- Minimal guarantees are still needed in order to reason about process execution
  - *weak fairness* is a useful abstract concept, every program is eventually scheduled to execute
  - the operating system scheduling policy needs to be assessed when making such an assumption

# Anomalies

## Atomicity

- Let  $x=3$  and  $y=5$
- Consider the statement  $x := x + y$
- What is the final value of  $x$ ?

# Anomalies

## Fairness

- Assume a priority-based non-preemptive schedule
- Process P has the high priority 1
- Process Q has the low priority 2
- P is idle
- Q is busy (running)
- When will P run again?



# Practical Concerns

In the absence of atomicity programming itself becomes impossible!

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance
  - add \$100
  - update balance
- Teller 2: deposit \$300
  - read account balance
  - add \$300
  - update balance
- Account balance

# Practical Concerns

In the absence of atomicity programming itself becomes impossible!

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance (\$245) (1)
  - add \$100
  - update balance
- Teller 2: deposit \$300
  - read account balance
  - add \$300
  - update balance
- Account balance

# Practical Concerns

In the absence of atomicity programming itself becomes impossible!

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance (\$245) (1)
  - add \$100 (\$345) (2)
  - update balance
- Teller 2: deposit \$300
  - read account balance
  - add \$300
  - update balance
- Account balance

# Practical Concerns

In the absence of atomicity programming itself becomes impossible!

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance (\$245) (1)
  - add \$100 (\$345) (2)
  - update balance
- Teller 2: deposit \$300
  - read account balance (\$245) (3)
  - add \$300
  - update balance
- Account balance

# Practical Concerns

In the absence of atomicity programming itself becomes impossible!

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance (\$245) (1)
  - add \$100 (\$345) (2)
  - update balance
- Teller 2: deposit \$300
  - read account balance (\$245) (3)
  - add \$300 (\$545) (4)
  - update balance
- Account balance

# Practical Concerns

In the absence of atomicity programming itself becomes impossible!

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance (\$245) (1)
  - add \$100 (\$345) (2)
  - update balance
- Teller 2: deposit \$300
  - read account balance (\$245) (3)
  - add \$300 (\$545) (4)
  - update balance (\$545) (5)
- Account balance

# Practical Concerns

In the absence of atomicity programming itself becomes impossible!

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance (\$245) (1)
  - add \$100 (\$345) (2)
  - update balance (\$345) (6)
- Teller 2: deposit \$300
  - read account balance (\$245) (3)
  - add \$300 (\$545) (4)
  - update balance (\$545) (5)
- Account balance **\$345 (WRONG!)**

# A Programming Language Solution

## Critical Region

- a block of code that is executed atomically
- a way to ensure mutual exclusion

Mutual Exclusion(mutex): Prevents simultaneous access to a shared resource



# A Programming Language Solution

This time, each deposit is a critical region.

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance
  - add \$100
  - update balance
- Teller 2: deposit \$300
  - read account balance
  - add \$300
  - update balance
- Account balance

# A Programming Language Solution

This time, each deposit is a critical region.

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance (\$245) (1)
  - add \$100 (\$345)
  - update balance (\$345)
- Teller 2: deposit \$300
  - read account balance
  - add \$300
  - update balance
- Account balance

# A Programming Language Solution

This time, each deposit is a critical region.

- Account balance \$245
- Teller 1: deposit \$100
  - read account balance (\$245) (1)
  - add \$100 (\$345)
  - update balance (\$345)
- Teller 2: deposit \$300
  - read account balance (\$345) (2)
  - add \$300 (\$645)
  - update balance (\$645)
- Account balance **\$645 (CORRECT!)**

# Basics of Mutual Exclusion

- Test and set
- Locks
- Semaphores
- Mutual exclusion constructs (programming language specific)

# Test and Set

Simple boolean flags cannot ensure mutual exclusion

- let  $G$  guard some resource that demands mutually exclusive access
- let  $G = \text{true}$  indicating that the resource is available
- processes  $P$  and  $Q$  need the resource

# Test and Set

Simple boolean flags cannot ensure mutual exclusion

- let  $G$  guard some resource that demands mutually exclusive access
- let  $G = \text{true}$  indicating that the resource is available
- processes  $P$  and  $Q$  need the resource
- $P$  reads  $G$  to be true

# Test and Set

Simple boolean flags cannot ensure mutual exclusion

- let G guard some resource that demands mutually exclusive access
- let  $G = \text{true}$  indicating that the resource is available
- processes P and Q need the resource
- P reads G to be true
- Q reads G to be true

# Test and Set

Simple boolean flags cannot ensure mutual exclusion

- let G guard some resource that demands mutually exclusive access
- let  $G = \text{true}$  indicating that the resource is available
- processes P and Q need the resource
- P reads G to be true
- Q reads G to be true
- P sets G to false
- P starts using the resource



# Test and Set

Simple boolean flags cannot ensure mutual exclusion

- let G guard some resource that demands mutually exclusive access
- let  $G = \text{true}$  indicating that the resource is available
- processes P and Q need the resource
- P reads G to be true
- Q reads G to be true
- P sets G to false
- P starts using the resource
- Q sets G to false
- Q starts using the resource

# Test and Set

- Hardware support is needed
- A process must test and set the flag in a single atomic step

```
while (true) do
  if G then G := false Must be atomic
    use resource
    G := true
    break
  fi
od
```

- The busy-wait is a real problem!

# Locks

- Test and set enables the introduction of locks
- Associate a lock with each resource
- Bracket the use of the resource with the operations
  - lock(G)
    - returns only when the lock is set
    - the process is suspended avoiding busy-wait
  - unlock(G)

# Locks

- A process may secure multiple resources as needed

lock(file1)

lock(file2)

transfer data from file1 to file2

unlock(file2)

unlock(file1)

- Possible anomalies:
  - accessing the resource without locking
  - failing to issue the unlock
  - deadlock

# Deadlock Avoidance

- Deadlock occurs when two processes are waiting on each other to release some resource
- One way of avoiding deadlock is for all the processes to lock resources in the same order

# Semaphores

- A semaphore is a construct designed to allow at most  $k$  processes get access to a given resource
- When  $k$  is 1, the semaphore becomes a basic lock (a binary semaphore)
- Traditionally the two operations over a semaphore are
  - $P(s)$  – tests for zero and decrements  $s$  by one, if greater than zero
  - $V(s)$  – increments  $s$  by one indicating the release of the resource
- All processes must follow the same protocol
  - $P(s)$  — guards entry to the resource use of resource
  - $V(s)$  — frees the resource